

radiologists and clinicians in the detection and characterization of breast lesions during ultrasound examinations.




Figure 6: An example of Samsung S-Detect-system interface [8]

This system has a lot of pros: it provides real-time feedback to the clinician during the examination, enabling immediate assessment and decision-making regarding lesion characterization and management. More than that, the software provides standardized reporting templates that facilitate structured documentation of lesion characteristics, including malignancy probability scores and recommended management options. This promotes consistency and completeness in reporting. S-Detect seamlessly integrates with the RS80A ultrasound system's workflow, allowing for efficient and streamlined use within clinical practice. It is user-friendly and does not significantly disrupt the examination process.

But S-Detect is primarily designed for breast lesion characterization and may not be suitable for other types of lesions or organs. Its utility is limited to breast ultrasound examinations and may not address the full spectrum of diagnostic challenges encountered in clinical practice. While S-Detect is user-friendly, clinicians may require some time to familiarize themselves with the software's features and functionality. Training and ongoing education may be necessary to optimize its use and interpretation. Moreover, The implementation of S-Detect may incur additional costs associated with software licensing, training, and maintenance. Clinics and healthcare facilities must consider the financial implications before adopting the technology.

There are another people who have been automating ultrasound diagnostics of the thyroid gland using artificial intelligence. For example in 2021 scientists from Romania published their article "Intelligent Diagnosis of Thyroid Ultrasound Imaging Using an Ensemble of Deep Learning Methods".

They developed a CNN-VGG ensemble fused from two models: a pre-trained fine-tuned model VGG-19 and an efficient lightweight CNN model. The proposed ensemble method proved to be an excellent and stable classifier with a good performance in terms of overall sen-

sitive predictive value (98.05%). [9]

Also there are scientists from China who published an "Artificial intelligence in thyroid ultrasound" article in 2023. Their research is more focused on the prevention and early detection of the thyroid cancer. They also used deep learning algorithms to achieve this goal. They tests different types of DL-based neural networks. [7]

The research of the above-mentioned scientists has been very successful. Their authors placed great emphasis on training the neural network to make diagnoses and look for pathology in ultrasound diagnostic images.

In the current work, a simpler and more global approach is considered: the neural network does not diagnose, but only assists the doctor. Their joint work makes it possible to minimize the errors of both the doctor and the software. The approach is described using the example of thyroid gland examination, but it can also be used in ultrasound diagnostics of other organs.

Also in this article, it is proposed to analyze not individual images, but a video recording of the entire research process.

Based on the approach described below, it is planned to develop a software product in the future and implement it into the work of a medical institution in a test mode.

Also, in the future, it is planned to develop the idea in such a way as to process not the final product of the work of some software: a visual representation of the ultrasound process, but the initial product, that is, ultrasonic signals. This will make the processing process faster.

V Proposed approach

Currently, artificial intelligence has been used in medicine for a long time. Integrating artificial intelligence into the ultrasound diagnostic process is not the easiest task. After all, software needs time to analyze.

Many studies allow to process the result later: for example, MRI, X-ray and others. But a standard ultrasound examination involves the interpretation of the result by a doctor right at the time of the study. In this regard, the quality of the study directly depends on the experience and attentiveness of the doctor.

To do intelligent processing of MRI results, it is enough to simply install the appropriate program on your computer. Because the MRI is first fully performed and then interpreted. And due to the fact that the ultrasound examination is simultaneously performed and interpreted, a third-party computer is rarely used by a doctor for it. But connecting the software directly to the ultrasound machine is almost impossible, for two reasons. Firstly, devices from different manufacturers with different software are used for diagnostics, which is written in lowlevel languages and can be difficult to integrate with other more modern technologies. Secondly, as mentioned earlier, the software needs time to process the data.

After numerous consultations with specialists in the medical field, analyzing the situation and finding the best way to introduce artificial intelligence into the ultrasound diagnostic process, it was decided to record the research process in a video format file. Then the data is transferred to the computer. The video is divided into frames of 0.5 seconds of research. It is this time interval that will allow not to process the same images several times, but at the same time not to miss important changes. The frames are then processed by a neural network. At the end of processing, the software generates its own, it will highlight a suspicious area and comment on it. In this case, the doctor can either ignore the prompts of artificial intelligence, if he has already paid attention to this pathology, or put a sensor and review the moment of interest again.

Training a neural network for the automated analysis of thyroid gland ultrasonography images involves several key steps.

The first step is to gather a large dataset of thyroid ultrasound images. These images should cover a wide range of thyroid conditions, including cysts, tumors, nodules, and other pathologies. The dataset should be diverse and representative of the population being analyzed.

Once the dataset is collected, it needs to be preprocessed to ensure consistency and quality. This may involve resizing the images, standardizing the brightness and contrast, and removing noise or artifacts. Each image in the dataset needs to be labeled with the corresponding thyroid pathology, such as cyst, tumor, or normal. This step is crucial for supervised learning, where the neural network learns from labeled examples.

Then it is need to choose an appropriate neural network architecture for the task. Convolutional Neural Networks (CNNs) are commonly used for image classification tasks due to their ability to capture spatial hierarchies in data.

A Convolutional Neural Network (CNN) is a type of deep learning algorithm specifically designed for processing and analyzing visual data, such as images. It is inspired by the structure and function of the human visual cortex and is well-suited for tasks such as image classification, object detection, and image segmentation. CNNs can not only classify images but also localize the regions within the image that contain abnormalities. This is crucial in medical imaging tasks, as it allows clinicians to pinpoint the location of cysts or tumors within the thyroid gland. CNNs can be trained to output bounding boxes or segmentation masks that delineate the boundaries of detected abnormalities.

CNNs consist of multiple layers, including convolutional layers, pooling layers, and fully connected layers.

Convolutional layers are the building blocks of CNNs. They apply convolution operations to input images, using learnable filters (also called kernels) to extract features such as edges, textures, and patterns. These filters slide




Figure 7: Overview and details of a convolutional neural network (CNN) architecture for image recognition [10]

over the input image, computing dot products between the filter weights and the input pixels to produce feature maps. Multiple filters are used in each convolutional layer to capture different features.

Pooling layers downsample the feature maps produced by the convolutional layers, reducing their spatial dimensions while retaining the most important information.

The most common type of pooling operation is max pooling, where the maximum value within each region of the feature map is retained, effectively reducing the size of the feature maps.

Fully connected layers, also known as dense layers, are traditional neural network layers where every neuron is connected to every neuron in the previous and subsequent layers. These layers are typically used at the end of the CNN to map the extracted features to the output classes or labels.

As for the size of the training sample, it depends on various factors such as the complexity of the task, the diversity of the dataset, and the chosen neural network architecture. In general, larger datasets tend to yield better-performing models, especially for deep learning tasks.

However, the minimum size of the training sample required for effective model training can vary significantly

depending on the specific problem being addressed. It is essential to strike a balance between dataset size, computational resources, and model performance when designing the training pipeline. In the case of medical imaging tasks like thyroid ultrasound analysis, larger datasets with thousands to tens of thousands of labeled images are typically required to train accurate and robust models.

NNs can be trained effectively even with limited labeled data by employing data augmentation techniques. These techniques involve applying transformations such as rotation, scaling, flipping, and cropping to the input images, thereby augmenting the training dataset and improving the model's generalization ability.

Pre-trained CNN models, which have been trained on large-scale datasets such as ImageNet, can be fine-tuned for medical image analysis tasks with relatively small datasets. By leveraging the feature representations learned from generic image data, transfer learning enables CNNs to achieve better performance and faster convergence when applied to medical image datasets, including thyroid ultrasound images.

CNNs can provide insights into the decision-making process by generating heatmaps or saliency maps that highlight the regions of the image that contribute most to the model's predictions. This interpretability is valuable for clinicians, as it helps them understand why a particular diagnosis or classification was made by the CNN.

In summary, Convolutional Neural Networks offer powerful capabilities for automatically detecting and localizing cysts, tumors, and other abnormalities on thyroid ultrasound images. By learning complex patterns and structures from labeled data, CNNs can assist radiologists and clinicians in diagnosing thyroid pathologies more accurately and efficiently, leading to improved patient outcomes.

The expected result of the implementation of the approach should be a web application with artificial intelligence inside. Between the desktop and the web application, the choice fell on the second option. This is due to the fact that the neural network is able to independently learn additionally in the course of its work. To do this, she needs to have access to the results of working with the application of other users. It is more convenient to do this in a web format. It is also necessary to be able to refine the application. By updating web applications, the added changes will quickly appear to all users, unlike the desktop application, where each user will have to update it.

V Proposed approach

However, there are some serious pitfalls here.

With the web approach, a single server will have access to all application data. This violates the privacy

policy and the protection of the user's personal data. The solution to this problem was found in having a separate server for each medical facility. And also not to transfer user data to the application. Based on the specifics of this software, it can process anonymous data and this will not affect the result.

The second difficulty encountered along the way is the presence of artifacts in the research process. The neural network must learn how to process them.

Artifacts in ultrasonic diagnostic imaging refer to misleading features or distortions present in the ultrasound image that do not accurately represent the anatomical structures being examined. These artifacts can arise due to various factors, including the properties of the ultrasound beam, the interaction of ultrasound waves with tissues, equipment settings, patient characteristics, and operator technique. Understanding and mitigating artifacts are essential for ensuring the accuracy and reliability of ultrasound diagnoses.

There are different types of artifacts. Reverberation Artifacts occurs when sound waves bounce back and forth between two strong reflectors, creating multiple, evenly spaced echoes on the image. It can give the appearance of additional structures or false boundaries within tissues.

Shadowing occurs when sound waves are attenuated by highly reflective or dense structures, resulting in a hypoechoic or anechoic region behind the structure. This can obscure underlying structures and limit visualization.

Edge artifacts occur at the interfaces between tissues with different acoustic properties. They manifest as bright or dark lines along tissue boundaries and can distort the appearance of adjacent structures.

Noise in ultrasound images can result from electronic interference, acoustic reverberations, or random fluctuations in signal intensity. It can degrade image quality and reduce diagnostic accuracy.

Motion artifacts occur when there is movement of the patient or probe during image acquisition. This can lead to blurring or ghosting of structures and compromise image clarity.

Teaching a CNN to process artifacts in ultrasound images is not an easy task. But there are some ways to overcome it.

Adversarial training involves training the CNN simultaneously with a generator network that generates realistic artifacts and a discriminator network that distinguishes between real images and artifacts. This helps the CNN learn to discriminate between artifacts and true structures.

Constructing a dataset that includes annotated examples of various artifacts encountered in clinical practice can facilitate CNN training. Annotating images to identify regions affected by artifacts allows the CNN to learn to ignore or compensate for them during analysis.